

# Trial 1

## Loading Modules and functions

```
In [ ]: from sklearn.naive_bayes import GaussianNB
from sklearn.metrics import classification_report, accuracy_score, confusion_matrix
from sklearn.model_selection import train_test_split
import numpy as np
import pandas as pd
```

## Loading Data

```
In [ ]: df = pd.read_csv("../processed.csv")
df
```

Out[ ]:

	Initial_Price	Final_Price	Win_Flag	Mac_Flag	Linux_Flag	Positive_Reviews	Negative_Reviews	Memory_MB	Storage_MB	target
0	52.0	52.0	True	True	False	57.0	7.0	1024	50	1
1	0.0	0.0	True	True	False	53.0	6.0	2048	3072	1
2	0.0	0.0	True	False	False	133.0	69.0	2048	100	0
3	530.0	530.0	True	False	False	22.0	9.0	2048	500	0
4	229.0	229.0	True	True	True	226.0	44.0	2048	1500	1
...	...	...	...	...	...	...	...	...	...	...
57467	85.0	85.0	True	False	False	0.0	4.0	4096	200	-1
57468	349.0	349.0	True	True	False	2.0	1.0	1024	1024	1
57469	164.0	164.0	True	False	False	8.0	1.0	4096	20480	1
57470	610.0	610.0	True	False	False	1.0	0.0	4096	3072	1
57471	570.0	285.0	True	False	False	0.0	1.0	1024	2048	-1

57472 rows × 10 columns

## Splitting Data 33% test, 66% train

```
In [ ]: y = df["target"]
X = df.drop(labels=["target"],axis=1)

X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.33, random_state=100,shuffle=True)
```

## Fitting model

```
In [ ]: bayes = GaussianNB()
bayes.fit(X_train, y_train)
```

Out[ ]:

GaussianNB ⓘ ?

GaussianNB()

## Predicting Likelihood of test cases

```
In [ ]: bayes.predict_proba(X_test)
```

```
Out[ ]: array([[3.77478260e-01, 6.22521738e-01, 1.64695980e-09],
 [3.77478260e-01, 6.22521738e-01, 1.64695980e-09],
 [3.77478260e-01, 6.22521738e-01, 1.64695980e-09],
 ...,
 [3.77478260e-01, 6.22521738e-01, 1.64695980e-09],
 [3.77478260e-01, 6.22521738e-01, 1.64695980e-09],
 [3.77478260e-01, 6.22521738e-01, 1.64695980e-09]])
```

## Model evaluation

```
In [ ]: preds = bayes.predict(X_test)
print(f"Accuracy : {accuracy_score(y_test, preds)}")
```

Accuracy : 0.18786249077296213

## Important metrics

```
In [ ]: print(classification_report(y_test, preds))
```

	precision	recall	f1-score	support
-1	0.00	0.00	0.00	2123
0	0.19	1.00	0.32	3563
1	0.00	0.00	0.00	13280
accuracy			0.19	18966
macro avg	0.06	0.33	0.11	18966
weighted avg	0.04	0.19	0.06	18966

c:\Users\user\AppData\Local\Programs\Python\Python311\Lib\site-packages\sklearn\metrics\\_classification.py:1509: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero\_division` parameter to control this behavior.  
\_warn\_prf(average, modifier, f"{metric.capitalize()} is", len(result))  
c:\Users\user\AppData\Local\Programs\Python\Python311\Lib\site-packages\sklearn\metrics\\_classification.py:1509: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero\_division` parameter to control this behavior.  
\_warn\_prf(average, modifier, f"{metric.capitalize()} is", len(result))  
c:\Users\user\AppData\Local\Programs\Python\Python311\Lib\site-packages\sklearn\metrics\\_classification.py:1509: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero\_division` parameter to control this behavior.  
\_warn\_prf(average, modifier, f"{metric.capitalize()} is", len(result))

## Confusion Matrix Display

```
In [ ]: confusion_matrix(y_test,preds)
```

Out[ ]: array([[ 0, 2123, 0],  
[ 0, 3563, 0],  
[ 0, 13280, 0]], dtype=int64)

## 1st trial Conclusion

- 1. Bayes Classifier gave *worst* performance of 18% accuracy
- 2. Its confusion matrix suggests classifier does no prediction of +1 and -1 class labels

# Trial 2 with normalised attributes

## Loading Modules and Functions

```
In [ ]: from sklearn.naive_bayes import GaussianNB
from sklearn.metrics import classification_report, accuracy_score, confusion_matrix
from sklearn.model_selection import train_test_split
import numpy as np
import pandas as pd
```

## Loading Data

```
In [ ]: df = pd.read_csv("../processed.csv")
df
```

Out[ ]:

	Initial_Price	Final_Price	Win_Flag	Mac_Flag	Linux_Flag	Positive_Reviews	Negative_Reviews	Memory_MB	Storage_MB	target
0	52.0	52.0	True	True	False	57.0	7.0	1024	50	1
1	0.0	0.0	True	True	False	53.0	6.0	2048	3072	1
2	0.0	0.0	True	False	False	133.0	69.0	2048	100	0
3	530.0	530.0	True	False	False	22.0	9.0	2048	500	0
4	229.0	229.0	True	True	True	226.0	44.0	2048	1500	1
...	...	...	...	...	...	...	...	...	...	...
57467	85.0	85.0	True	False	False	0.0	4.0	4096	200	-1
57468	349.0	349.0	True	True	False	2.0	1.0	1024	1024	1
57469	164.0	164.0	True	False	False	8.0	1.0	4096	20480	1
57470	610.0	610.0	True	False	False	1.0	0.0	4096	3072	1
57471	570.0	285.0	True	False	False	0.0	1.0	1024	2048	-1

57472 rows × 10 columns

## Normalising Continous Features

```
In [ ]: def normalise(feature, df):
    mean = df[feature].mean()
    sd = df[feature].std()
    df[feature] = (df[feature] - mean) / sd

normalise("Initial_Price", df)
normalise("Final_Price", df)
normalise("Positive_Reviews", df)
normalise("Negative_Reviews", df)
normalise("Memory_MB", df)
normalise("Storage_MB", df)
df
```

Out[ ]:

	Initial_Price	Final_Price	Win_Flag	Mac_Flag	Linux_Flag	Positive_Reviews	Negative_Reviews	Memory_MB	Storage_MB	target
0	-0.271301	-0.258274	True	True	False	-0.034488	-0.033031	-0.004171	-0.004171	1
1	-0.322070	-0.309572	True	True	False	-0.034609	-0.033196	-0.004171	-0.004171	1
2	-0.322070	-0.309572	True	False	False	-0.032193	-0.022842	-0.004171	-0.004171	0
3	0.195385	0.213271	True	False	False	-0.035545	-0.032703	-0.004171	-0.004171	0
4	-0.098490	-0.083665	True	True	True	-0.029385	-0.026950	-0.004171	-0.004171	1
...	...	...	...	...	...	...	...	...	...	...
57467	-0.239082	-0.225720	True	False	False	-0.036210	-0.033524	-0.004171	-0.004171	-1
57468	0.018669	0.034715	True	True	False	-0.036149	-0.034017	-0.004171	-0.004171	1
57469	-0.161952	-0.147787	True	False	False	-0.035968	-0.034017	-0.004171	-0.004171	1
57470	0.273492	0.292190	True	False	False	-0.036179	-0.034182	-0.004171	-0.004171	1
57471	0.234439	-0.028421	True	False	False	-0.036210	-0.034017	-0.004171	-0.004171	-1

57472 rows × 10 columns

## Splitting data to 33% Test , 66% Train

```
In [ ]: y = df["target"]
X = df.drop(labels=["target","Win_Flag","Mac_Flag","Linux_Flag"],axis=1)

X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.33, random_state=100,shuffle=True)
```

## Fitting the model

```
In [ ]: bayes = GaussianNB()
bayes.fit(X_train, y_train)
```

```
Out[ ]: GaussianNB
GaussianNB()
```

## Likelihood of predictions

```
In [ ]: bayes.predict_proba(X_test)
```

```
Out[ ]: array([[9.90667437e-01, 9.33256327e-03, 9.81563140e-12],
               [9.91290910e-01, 8.70909043e-03, 9.19860142e-12],
               [9.91237383e-01, 8.76261690e-03, 9.22623962e-12],
               ...,
               [9.91328087e-01, 8.67191274e-03, 9.17865198e-12],
               [9.87943714e-01, 1.20562857e-02, 1.28388006e-11],
               [9.91392419e-01, 8.60758149e-03, 9.16860493e-12]])
```

## Evaluating Model with test data

```
In [ ]: preds = bayes.predict(X_test)
print(f"Accuracy : {accuracy_score(y_test, preds)}")
```

Accuracy : 0.11879152167035748

## Important Metrics

```
In [ ]: print(classification_report(y_test, preds))
```

	precision	recall	f1-score	support
-1	0.12	1.00	0.21	2123
0	0.10	0.03	0.05	3563
1	1.00	0.00	0.00	13280
accuracy			0.12	18966
macro avg	0.41	0.34	0.09	18966
weighted avg	0.73	0.12	0.04	18966

## Display Confusion Matrix

```
In [ ]: confusion_matrix(y_test,preds)
```

```
Out[ ]: array([[ 2116,    7,    0],
               [ 3451,  112,    0],
               [12224, 1031,   25]], dtype=int64)
```

## 2nd trial Conclusion

1. Bayes Classifier's performance fall to 11% accuracy , after normalising attributes and discarding flags to keep only continous attributes
2. This follows that the continous features of dataset were not *guassian* , that is they are not normally distributed.